

```

1 struct sock {
2     struct sock_common    __sk_common; /* early demux fields */
3     struct dst_entry __rcu *sk_rx_dst;
4     int                   sk_rx_dst_ifindex;
5     u32                   sk_rx_dst_cookie;
6     socket_lock_t         sk_lock;
7     atomic_t              sk_drops;
8     int                   sk_rcvlowat;
9     struct sk_buff_head    sk_error_queue;
10    struct sk_buff_head    sk_receive_queue;
11    struct {
12        atomic_t           rmem_alloc;
13        int                len;
14        struct sk_buff      *head;
15        struct sk_buff      *tail;
16    } sk_backlog;
17    int                    sk_forward_alloc;
18    u32                    sk_reserved_mem;
19    #ifdef CONFIG_NET_RX_BUSY_POLL
20    unsigned int           sk_ll_usec;
21    /* ===== mostly read cache line ===== */
22    unsigned int           sk_napi_id;
23    #endif
24    int                    sk_rcvbuf;
25    int                    sk_disconnects;
26    struct sk_filter __rcu *sk_filter;
27    union {
28        struct socket_wq __rcu *sk_wq;
29        /* private: */
30        struct socket_wq      *sk_wq_raw;
31        /* public: */
32    };
33    #ifdef CONFIG_XFRM
34    struct xfrm_policy __rcu *sk_policy[2];
35    #endif
36    struct dst_entry __rcu *sk_dst_cache;
37    atomic_t           sk_omem_alloc;
38    int                sk_sndbuf;
39    /* ===== cache line for TX ===== */
40    int                sk_wmem_queued;
41    refcount_t         sk_wmem_alloc;
42    unsigned long       sk_tsq_flags;
43    union {
44        struct sk_buff *sk_send_head;
45        struct rb_root tcp_rtx_queue;
46    };
47    struct sk_buff_head sk_write_queue;
48    __s32               sk_peek_off;
49    int                 sk_write_pending;
50    __u32               sk_dst_pending_confirm;
51    u32                 sk_pacing_status; /* see enum sk_pacing */
52    long                sk_sndtimeo;
53    struct timer_list    sk_timer;
54    __u32               sk_priority;
55    __u32               sk_mark;
56    unsigned long        sk_pacing_rate; /* bytes per second */
57    unsigned long        sk_max_pacing_rate;
58    struct page_frag     sk_frag;
59    netdev_features_t     sk_route_caps;
60    int                  sk_gso_type;
61    unsigned int          sk_gso_max_size;
62    gfp_t                sk_allocation;
63    __u32                sk_txhash;
64    /*
65     * Because of non atomicity rules, all
66     * changes are protected by socket lock.
67     */
68    u8                   sk_gso_disabled : 1,
69    sk_kern_sock : 1,
70    sk_no_check_tx : 1,

```

```

71     sk_no_check_rx : 1,
72     sk_userlocks : 4;
73     u8                sk_pacing_shift;
74     u16               sk_type;
75     u16               sk_protocol;
76     u16               sk_gso_max_segs;
77     unsigned long     sk_lingertime;
78     struct proto      *sk_prot_creator;
79     rwlock_t          sk_callback_lock;
80     int               sk_err,
81     sk_err_soft;
82     u32               sk_ack_backlog;
83     u32               sk_max_ack_backlog;
84     kuid_t            sk_uid;
85     u8                sk_txrehash;
86     #ifdef CONFIG_NET_RX_BUSY_POLL
87     u8                sk_prefer_busy_poll;
88     u16               sk_busy_poll_budget;
89     #endif
90     spinlock_t         sk_peer_lock;
91     int               sk_bind_phc;
92     struct pid         *sk_peer_pid;
93     const struct cred  *sk_peer_cred;
94     long               sk_rcvtimeo;
95     ktime_t            sk_stamp;
96     #if BITS_PER_LONG==32
97     seqlock_t          sk_stamp_seq;
98     #endif
99     atomic_t           sk_tskey;
100    atomic_t           sk_zckey;
101    u32               sk_tsflags;
102    u8                sk_shutdown;
103    u8                sk_clockid;
104    u8                sk_txtime_deadline_mode : 1,
105    sk_txtime_report_errors : 1,
106    sk_txtime_unused : 6;
107    bool               sk_use_task_frag;
108    struct socket       *sk_socket;
109    void               *sk_user_data;
110    #ifdef CONFIG_SECURITY
111    void               *sk_security;
112    #endif
113    struct sock_cgroup_data sk_cgrp_data;
114    struct mem_cgroup   *sk_memcg;
115    void               (*sk_state_change)(struct sock *sk);
116    void               (*sk_data_ready)(struct sock *sk);
117    void               (*sk_write_space)(struct sock *sk);
118    void               (*sk_error_report)(struct sock *sk);
119    int                (*sk_backlog_rcv)(struct sock *sk,
120    struct sk_buff *skb);
121    #ifdef CONFIG_SOCK_VALIDATE_XMIT
122    struct sk_buff*     (*sk_validate_xmit_skb)(struct sock *sk,
123    struct net_device *dev,
124    struct sk_buff *skb);
125    #endif
126    void               (*sk_destruct)(struct sock *sk);
127    struct sock_reuseport __rcu *sk_reuseport_cb;
128    #ifdef CONFIG_BPF_SYSCALL
129    struct bpf_local_storage __rcu *sk_bpf_storage;
130    #endif
131    struct rcu_head     sk_rcu;
132    netns_tracker       ns_tracker;
133 };

```

```

1 struct socket {
2     socket_state      state;
3     short             type;
4     unsigned long     flags;
5     struct file       *file;
6     struct sock       *sk;
7     const struct proto_ops *ops; /* Might change with IPV6_ADDRFORM or MPTCP. */
8     struct socket_wq  wq;
9 };

```

```

1 struct sockaddr {
2     sa_family_t      sa_family; /* address family, AF_xxx */
3     char sa_data_min[14];
4 };
5
6 struct in_addr {
7     u32 s_addr;
8 };
9
10 struct sockaddr_in {
11     __kernel_sa_family_t sin_family; /* Address family */
12     __be16 sin_port; /* Port number */
13     struct in_addr sin_addr; /* Internet address*/
14
15     /* Pad to size of 'struct sockaddr'. */
16     unsigned char __pad[__SOCK_SIZE__ - sizeof(short int) - sizeof(unsigned short int) -
17     sizeof(struct in_addr)];
17 };

```

```

1 struct epoll_event {
2     __poll_t events;
3     __u64 data;
4 };

```

```

1 struct super_block {
2     struct list_head      s_list;           /* Keep this first */
3     dev_t                 s_dev;           /* search index; _not_ kdev_t */
4     unsigned char         s_blocksize_bits;
5     unsigned long         s_blocksize;
6     loff_t                s_maxbytes;      /* Max file size */
7     struct file_system_type *s_type;
8     const struct super_operations *s_op;
9     const struct dquot_operations *dq_op;
10    const struct quotactl_ops *s_qcop;
11    const struct export_operations *s_export_op;
12    unsigned long          s_flags;
13    unsigned long          s_iflags;        /* internal SB_I_* flags */
14    unsigned long          s_magic;
15    struct dentry          *s_root;
16    struct rw_semaphore    s_umount;
17    int                    s_count;
18    atomic_t               s_active;
19    #ifdef CONFIG_SECURITY
20    void                   *s_security;
21    #endif
22    const struct xattr_handler * const *s_xattr;
23    #ifdef CONFIG_FS_ENCRYPTION
24    const struct fscrypt_operations *s_cop;
25    struct fscrypt_keyring *s_master_keys; /* master crypto keys in use */
26    #endif
27    #ifdef CONFIG_FS_VERITY
28    const struct fsverity_operations *s_vop;
29    #endif
30    #if IS_ENABLED(CONFIG_UNICODE)
31    struct unicode_map *s_encoding;
32    __u16 s_encoding_flags;
33    #endif
34    struct hlist_bl_head    s_roots;         /* alternate root dentries for NFS */
35    struct list_head        s_mounts;        /* list of mounts; _not_ for fs use */
36    struct block_device     *s_bdev;
37    struct bdev_handle      *s_bdev_handle;
38    struct backing_dev_info *s_bdi;
39    struct mtd_info         *s_mtd;
40    struct hlist_node       s_instances;
41    unsigned int            s_quota_types;    /* Bitmask of supported quota types */
42    struct quota_info        s_dquot;        /* Diskquota specific options */
43
44    struct sb_writers        s_writers;
45
46    /*
47     * Keep s_fs_info, s_time_gran, s_fsnotify_mask, and
48     * s_fsnotify_marks together for cache efficiency. They are frequently
49     * accessed and rarely modified.
50     */
51    void                   *s_fs_info;      /* Filesystem private info */
52
53    /* Granularity of c/m/ctime in ns (cannot be worse than a second) */
54    u32                    s_time_gran;
55    /* Time limits for c/m/ctime in seconds */
56    time64_t               s_time_min;
57    time64_t               s_time_max;
58    #ifdef CONFIG_FSNOTIFY
59    __u32                   s_fsnotify_mask;
60    struct fsnotify_mark_connector __rcu *s_fsnotify_marks;
61    #endif
62
63    char                   s_id[32];        /* Informational name */
64    uuid_t                 s_uuid;          /* UUID */
65
66    unsigned int            s_max_links;
67
68    /*
69     * The next field is for VFS *only*. No filesystems have any business
70     * even looking at it. You had been warned.

```

```

71  */
72  struct mutex s_vfs_rename_mutex;          /* Kludge */
73
74  /*
75  * Filesystem subtype.  If non-empty the filesystem type field
76  * in /proc/mounts will be "type.subtype"
77  */
78  const char *s_subtype;
79
80  const struct dentry_operations *s_d_op; /* default d_op for dentries */
81
82  struct shrinker *s_shrink;                /* per-sb shrinker handle */
83
84  /* Number of inodes with nlink == 0 but still referenced */
85  atomic_long_t s_remove_count;
86
87  /*
88  * Number of inode/mount/sb objects that are being watched, note that
89  * inodes objects are currently double-accounted.
90  */
91  atomic_long_t s_fsnotify_connectors;
92
93  /* Read-only state of the superblock is being changed */
94  int s_readonly_remount;
95
96  /* per-sb errseq_t for reporting writeback errors via syncfs */
97  errseq_t s_wb_err;
98
99  /* AIO completions deferred from interrupt context */
100 struct workqueue_struct *s_dio_done_wq;
101 struct hlist_head s_pins;
102
103 /*
104 * Owning user namespace and default context in which to
105 * interpret filesystem uids, gids, quotas, device nodes,
106 * xattrs and security labels.
107 */
108 struct user_namespace *s_user_ns;
109
110 /*
111 * The list_lru structure is essentially just a pointer to a table
112 * of per-node lru lists, each of which has its own spinlock.
113 * There is no need to put them into separate cachelines.
114 */
115 struct list_lru          s_dentry_lru;
116 struct list_lru          s_inode_lru;
117 struct rcu_head          rcu;
118 struct work_struct       destroy_work;
119
120 struct mutex             s_sync_lock;      /* sync serialisation lock */
121
122 /*
123 * Indicates how deep in a filesystem stack this SB is
124 */
125 int s_stack_depth;
126
127 /* s_inode_list_lock protects s_inodes */
128 spinlock_t              s_inode_list_lock ____cacheline_aligned_in_smp;
129 struct list_head         s_inodes;        /* all inodes */
130
131 spinlock_t              s_inode_wblist_lock;
132 struct list_head         s_inodes_wb;     /* writeback inodes */
133 } __randomize_layout;

```

```

1 struct super_operations {
2     struct inode *(*alloc_inode)(struct super_block *sb);
3     void (*destroy_inode)(struct inode *);
4     void (*free_inode)(struct inode *);
5
6     void (*dirty_inode) (struct inode *, int flags);
7     int (*write_inode) (struct inode *, struct writeback_control *wbc);
8     int (*drop_inode) (struct inode *);
9     void (*evict_inode) (struct inode *);
10    void (*put_super) (struct super_block *);
11    int (*sync_fs)(struct super_block *sb, int wait);
12    int (*freeze_super) (struct super_block *, enum freeze_holder who);
13    int (*freeze_fs) (struct super_block *);
14    int (*thaw_super) (struct super_block *, enum freeze_holder who);
15    int (*unfreeze_fs) (struct super_block *);
16    int (*statfs) (struct dentry *, struct kstatfs *);
17    int (*remount_fs) (struct super_block *, int *, char *);
18    void (*umount_begin) (struct super_block *);
19
20    int (*show_options)(struct seq_file *, struct dentry *);
21    int (*show_devname)(struct seq_file *, struct dentry *);
22    int (*show_path)(struct seq_file *, struct dentry *);
23    int (*show_stats)(struct seq_file *, struct dentry *);
24    #ifdef CONFIG_QUOTA
25    ssize_t (*quota_read)(struct super_block *, int, char *, size_t, loff_t);
26    ssize_t (*quota_write)(struct super_block *, int, const char *, size_t, loff_t);
27    struct dquot __rcu **(*get_dquots)(struct inode *);
28    #endif
29    long (*nr_cached_objects)(struct super_block *,
30    struct shrink_control *);
31    long (*free_cached_objects)(struct super_block *,
32    struct shrink_control *);
33    void (*shutdown)(struct super_block *sb);
34 };

```

```

1 struct dentry {
2     /* RCU lookup touched fields */
3     unsigned int d_flags;           /* protected by d_lock */
4     seqcount_spinlock_t d_seq;     /* per dentry seqlock */
5     struct hlist_bl_node d_hash;    /* lookup hash list */
6     struct dentry *d_parent;        /* parent directory */
7     struct qstr d_name;
8     struct inode *d_inode;          /* Where the name belongs to - NULL is
9     * negative */
10    unsigned char d_iname[DNAME_INLINE_LEN];    /* small names */
11
12    /* Ref lookup also touches following */
13    struct lockref d_lockref;        /* per-dentry lock and refcount */
14    const struct dentry_operations *d_op;
15    struct super_block *d_sb;        /* The root of the dentry tree */
16    unsigned long d_time;            /* used by d_revalidate */
17    void *d_fsdata;                 /* fs-specific data */
18
19    union {
20        struct list_head d_lru;      /* LRU list */
21        wait_queue_head_t *d_wait;   /* in-lookup ones only */
22    };
23    struct hlist_node d_sib;          /* child of parent list */
24    struct hlist_head d_children;     /* our children */
25    /*
26    * d_alias and d_rcu can share memory
27    */
28    union {
29        struct hlist_node d_alias;    /* inode alias list */
30        struct hlist_bl_node d_in_lookup_hash; /* only for in-lookup ones */
31        struct rcu_head d_rcu;
32    } d_u;
33 };

```

```

1 struct dentry_operations {
2     int (*d_revalidate)(struct dentry *, unsigned int);
3     int (*d_weak_revalidate)(struct dentry *, unsigned int);
4     int (*d_hash)(const struct dentry *, struct qstr *);
5     int (*d_compare)(const struct dentry *,
6     unsigned int, const char *, const struct qstr *);
7     int (*d_delete)(const struct dentry *);
8     int (*d_init)(struct dentry *);
9     void (*d_release)(struct dentry *);
10    void (*d_prune)(struct dentry *);
11    void (*d_iput)(struct dentry *, struct inode *);
12    char *(*d_dname)(struct dentry *, char *, int);
13    struct vfsmount *(*d_automount)(struct path *);
14    int (*d_manage)(const struct path *, bool);
15    struct dentry *(*d_real)(struct dentry *, const struct inode *);
16 } ____cacheline_aligned;

```

```

1 struct inode {
2     umode_t                i_mode;
3     unsigned short         i_opflags;
4     kuid_t                 i_uid;
5     kgid_t                 i_gid;
6     unsigned int           i_flags;
7     #ifdef CONFIG_FS_POSIX_ACL
8     struct posix_acl        *i_acl;
9     struct posix_acl        *i_default_acl;
10    #endif
11    const struct inode_operations *i_op;
12    struct super_block        *i_sb;
13    struct address_space      *i_mapping;
14    #ifdef CONFIG_SECURITY
15    void                      *i_security;
16    #endif
17    /* Stat data, not accessed from path walking */
18    unsigned long            i_ino;
19    /* Filesystems may only read i_nlink directly. */
20    union {
21        const unsigned int i_nlink;
22        unsigned int __i_nlink;
23    };
24    dev_t                    i_rdev;
25    loff_t                   i_size;
26    struct timespec64        __i_atime;
27    struct timespec64        __i_mtime;
28    struct timespec64        __i_ctime; /* use inode*_ctime accessors! */
29    spinlock_t               i_lock; /* i_blocks, i_bytes, maybe i_size */
30    unsigned short           i_bytes;
31    u8                       i_blkbits;
32    u8                       i_write_hint;
33    blkcnt_t                 i_blocks;
34    #ifdef __NEED_I_SIZE_ORDERED
35    seqcount_t               i_size_seqcount;
36    #endif
37    /* Misc */
38    unsigned long            i_state;
39    struct rw_semaphore       i_rwsem;
40    unsigned long            dirtied_when; /* jiffies of first dirtying */
41    unsigned long            dirtied_time_when;
42    struct hlist_node         i_hash;
43    struct list_head          i_io_list; /* backing dev IO list */
44    #ifdef CONFIG_CGROUP_WRITEBACK
45    struct bdi_writeback      *i_wb; /* the associated cgroup wb */
46    /* foreign inode detection, see wbc_detach_inode() */
47    int                       i_wb_frn_winner;
48    u16                      i_wb_frn_avg_time;
49    u16                      i_wb_frn_history;
50    #endif
51    struct list_head          i_lru; /* inode LRU list */
52    struct list_head          i_sb_list;
53    struct list_head          i_wb_list; /* backing dev writeback list */
54    union {
55        struct hlist_head i_dentry;
56        struct rcu_head i_rcu;
57    };
58    atomic64_t                i_version;
59    atomic64_t                i_sequence; /* see futex */
60    atomic_t                  i_count;
61    atomic_t                  i_dio_count;
62    atomic_t                  i_writecount;
63    #if defined(CONFIG_IMA) || defined(CONFIG_FILE_LOCKING)
64    atomic_t                  i_readcount; /* struct files open RO */
65    #endif
66    union {
67        const struct file_operations *i_fop; /* former ->i_op->default_file_ops */
68        void (*free_inode)(struct inode *);
69    };
70    struct file_lock_context  *i_flctx;

```



```

71 struct address_space      i_data;
72 struct list_head          i_devices;
73 union {
74     struct pipe_inode_info *i_pipe;
75     struct cdev            *i_cdev;
76     char                   *i_link;
77     unsigned               i_dir_seq;
78 };
79 __u32                      i_generation;
80 #ifdef CONFIG_FSNOTIFY
81 __u32                      i_fsnotify_mask; /* all events this inode cares about */
82 struct fsnotify_mark_connector __rcu    *i_fsnotify_marks;
83 #endif
84 #ifdef CONFIG_FS_ENCRYPTION
85 struct fscrypt_inode_info    *i_crypt_info;
86 #endif
87 #ifdef CONFIG_FS_VERITY
88 struct fsverity_info         *i_verity_info;
89 #endif
90 void                        *i_private; /* fs or device private pointer */
91 } __randomize_layout;

```

```

1 struct inode_operations {
2     struct dentry * (*lookup) (struct inode *, struct dentry *, unsigned int);
3     const char * (*get_link) (struct dentry *, struct inode *, struct delayed_call *);
4     int (*permission) (struct mnt_idmap *, struct inode *, int);
5     struct posix_acl * (*get_inode_acl) (struct inode *, int, bool);
6
7     int (*readlink) (struct dentry *, char __user *, int);
8
9     int (*create) (struct mnt_idmap *, struct inode *, struct dentry *,
10     umode_t, bool);
11     int (*link) (struct dentry *, struct inode *, struct dentry *);
12     int (*unlink) (struct inode *, struct dentry *);
13     int (*symlink) (struct mnt_idmap *, struct inode *, struct dentry *,
14     const char *);
15     int (*mkdir) (struct mnt_idmap *, struct inode *, struct dentry *,
16     umode_t);
17     int (*rmdir) (struct inode *, struct dentry *);
18     int (*mknod) (struct mnt_idmap *, struct inode *, struct dentry *,
19     umode_t, dev_t);
20     int (*rename) (struct mnt_idmap *, struct inode *, struct dentry *,
21     struct inode *, struct dentry *, unsigned int);
22     int (*setattr) (struct mnt_idmap *, struct dentry *, struct iattr *);
23     int (*getattr) (struct mnt_idmap *, const struct path *,
24     struct kstat *, u32, unsigned int);
25     ssize_t (*listxattr) (struct dentry *, char *, size_t);
26     int (*fiemap) (struct inode *, struct fiemap_extent_info *, u64 start,
27     u64 len);
28     int (*update_time) (struct inode *, int);
29     int (*atomic_open) (struct inode *, struct dentry *,
30     struct file *, unsigned open_flag,
31     umode_t create_mode);
32     int (*tmpfile) (struct mnt_idmap *, struct inode *,
33     struct file *, umode_t);
34     struct posix_acl * (*get_acl) (struct mnt_idmap *, struct dentry *,
35     int);
36     int (*set_acl) (struct mnt_idmap *, struct dentry *,
37     struct posix_acl *, int);
38     int (*fileattr_set) (struct mnt_idmap *idmap,
39     struct dentry *dentry, struct fileattr *fa);
40     int (*fileattr_get) (struct dentry *dentry, struct fileattr *fa);
41     struct offset_ctx * (*get_offset_ctx) (struct inode *inode);
42 } ____cacheline_aligned;

```

```

1 struct file_system_type {
2     const char *name;
3     int fs_flags;
4     #define FS_REQUIRES_DEV      1
5     #define FS_BINARY_MOUNTDATA  2
6     #define FS_HAS_SUBTYPE       4
7     #define FS_USERNS_MOUNT      8      /* Can be mounted by usersns root */
8     #define FS_DISALLOW_NOTIFY_PERM 16   /* Disable fanotify permission events */
9     #define FS_ALLOW_IDMAP        32     /* FS has been updated to handle vfs
idmappings. */
10    #define FS_RENAME_DOES_D_MOVE 32768  /* FS will handle d_move() during rename()
internally. */
11    int (*init_fs_context)(struct fs_context *);
12    const struct fs_parameter_spec *parameters;
13    struct dentry *(*mount)(struct file_system_type *, int,
14    const char *, void *);
15    void (*kill_sb)(struct super_block *);
16    struct module *owner;
17    struct file_system_type *next;
18    struct hlist_head fs_supers;
19
20    struct lock_class_key s_lock_key;
21    struct lock_class_key s_umount_key;
22    struct lock_class_key s_vfs_rename_key;
23    struct lock_class_key s_writers_key[SB_FREEZE_LEVELS];
24
25    struct lock_class_key i_lock_key;
26    struct lock_class_key i_mutex_key;
27    struct lock_class_key invalidate_lock_key;
28    struct lock_class_key i_mutex_dir_key;
29 };

```

```

1 struct fs_struct {
2     int users;
3     spinlock_t lock;
4     seqcount_spinlock_t seq;
5     int umask;
6     int in_exec;
7     struct path root, pwd;
8 } __randomize_layout;

```

```

1 struct path {
2     struct vfsmount *mnt;
3     struct dentry *dentry;
4 } __randomize_layout;

```

```

1 struct files_struct {
2     /*
3      * read mostly part
4      */
5     atomic_t count;
6     bool resize_in_progress;
7     wait_queue_head_t resize_wait;
8     struct fdtable __rcu *fdt;
9     struct fdtable fdtab;
10    /*
11     * written part on a separate cache line in SMP
12     */
13    spinlock_t file_lock ____cacheline_aligned_in_smp;
14    unsigned int next_fd;
15    unsigned long close_on_exec_init[1];
16    unsigned long open_fds_init[1];
17    unsigned long full_fds_bits_init[1];
18    struct file __rcu *fd_array[NR_OPEN_DEFAULT];
19 };

```

```

1 struct file {
2     union {
3         /* fput() uses task work when closing and freeing file (default). */
4         struct callback_head    f_task_work;
5         /* fput() must use workqueue (most kernel threads). */
6         struct llist_node       f_llist;
7         unsigned int            f_iocb_flags;
8     };
9     /* Protects f_ep, f_flags. Must not be taken from IRQ context. */
10    spinlock_t                   f_lock;
11    fmode_t                      f_mode;
12    atomic_long_t                f_count;
13    struct mutex                 f_pos_lock;
14    loff_t                      f_pos;
15    unsigned int                 f_flags;
16    struct fown_struct           f_owner;
17    const struct cred            *f_cred;
18    struct file_ra_state         f_ra;
19    struct path                  f_path;
20    struct inode                 *f_inode;          /* cached value */
21    const struct file_operations *f_op;
22    u64                          f_version;
23    #ifdef CONFIG_SECURITY
24    void                         *f_security;
25    #endif
26    /* needed for tty driver, and maybe others */
27    void                         *private_data;
28    #ifdef CONFIG_EPOLL
29    /* Used by fs/eventpoll.c to link all the hooks to this file */
30    struct hlist_head            *f_ep;
31    #endif /* #ifdef CONFIG_EPOLL */
32    struct address_space         *f_mapping;
33    errseq_t                     f_wb_err;
34    errseq_t                     f_sb_err; /* for syncfs */
35 } __randomize_layout
36 __attribute__((aligned(4)));

```

```

1 struct _IO_FILE
2 {
3     int _flags; /* High-order word is _IO_MAGIC; rest is flags. */
4     /* The following pointers correspond to the C++ streambuf protocol. */
5     char *_IO_read_ptr; /* Current read pointer */
6     char *_IO_read_end; /* End of get area. */
7     char *_IO_read_base; /* Start of putback+get area. */
8     char *_IO_write_base; /* Start of put area. */
9     char *_IO_write_ptr; /* Current put pointer. */
10    char *_IO_write_end; /* End of put area. */
11    char *_IO_buf_base; /* Start of reserve area. */
12    char *_IO_buf_end; /* End of reserve area. */
13    /* The following fields are used to support backing up and undo. */
14    char *_IO_save_base; /* Pointer to start of non-current get area. */
15    char *_IO_backup_base; /* Pointer to first valid character of backup area */
16    char *_IO_save_end; /* Pointer to end of non-current get area. */
17    struct _IO_marker *_markers;
18
19    struct _IO_FILE *_chain;
20    int _fileno;
21    int _flags2;
22    __off_t _old_offset; /* This used to be _offset but it's too small. */
23    /* 1+column number of pbase(); 0 is unknown. */
24    unsigned short _cur_column;
25    signed char _vtable_offset;
26    char _shortbuf[1];
27    _IO_lock_t *_lock;
28    #ifdef _IO_USE_OLD_IO_FILE
29 };

```

```

1 struct proc_dir_entry {
2     /*
3      * number of callers into module in progress;
4      * negative -> it's going away RSN
5      */
6     atomic_t in_use;
7     refcount_t refcnt;
8     struct list_head pde_openers; /* who did ->open, but not ->release */
9     /* protects ->pde_openers and all struct pde_opener instances */
10    spinlock_t pde_unload_lock;
11    struct completion *pde_unload_completion;
12    const struct inode_operations *proc_iops;
13    union {
14        const struct proc_ops *proc_ops;
15        const struct file_operations *proc_dir_ops;
16    };
17    const struct dentry_operations *proc_dops;
18    union {
19        const struct seq_operations *seq_ops;
20        int (*single_show)(struct seq_file *, void *);
21    };
22    proc_write_t write;
23    void *data;
24    unsigned int state_size;
25    unsigned int low_ino;
26    nlink_t nlink;
27    kuid_t uid;
28    kgid_t gid;
29    loff_t size;
30    struct proc_dir_entry *parent;
31    struct rb_root subdir;
32    struct rb_node subdir_node;
33    char *name;
34    umode_t mode;
35    u8 flags;
36    u8 namelen;
37    char inline_name[];
38 } __randomize_layout;

```

```

1 struct proc_ops {
2     unsigned int proc_flags;
3     int (*proc_open)(struct inode *, struct file *);
4     ssize_t (*proc_read)(struct file *, char __user *, size_t, loff_t *);
5     ssize_t (*proc_read_iter)(struct kiocb *, struct iov_iter *);
6     ssize_t (*proc_write)(struct file *, const char __user *, size_t, loff_t *);
7     /* mandatory unless nonseekable_open() or equivalent is used */
8     loff_t (*proc_lseek)(struct file *, loff_t, int);
9     int (*proc_release)(struct inode *, struct file *);
10    __poll_t (*proc_poll)(struct file *, struct poll_table_struct *);
11    long (*proc_ioctl)(struct file *, unsigned int, unsigned long);
12    #ifdef CONFIG_COMPAT
13    long (*proc_compat_ioctl)(struct file *, unsigned int, unsigned long);
14    #endif
15    int (*proc_mmap)(struct file *, struct vm_area_struct *);
16    unsigned long (*proc_get_unmapped_area)(struct file *, unsigned long, unsigned long,
17    unsigned long, unsigned long);
18 } __randomize_layout;

```

```

1 struct file_operations {
2     struct module *owner;
3     loff_t (*llseek) (struct file *, loff_t, int);
4     ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
5     ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
6     ssize_t (*read_iter) (struct kiocb *, struct iov_iter *);
7     ssize_t (*write_iter) (struct kiocb *, struct iov_iter *);
8     int (*iopoll)(struct kiocb *kiocb, struct io_comp_batch *,
9     unsigned int flags);
10    int (*iterate_shared) (struct file *, struct dir_context *);
11    __poll_t (*poll) (struct file *, struct poll_table_struct *);
12    long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long);
13    long (*compat_ioctl) (struct file *, unsigned int, unsigned long);
14    int (*mmap) (struct file *, struct vm_area_struct *);
15    unsigned long mmap_supported_flags;
16    int (*open) (struct inode *, struct file *);
17    int (*flush) (struct file *, fl_owner_t id);
18    int (*release) (struct inode *, struct file *);
19    int (*fsync) (struct file *, loff_t, loff_t, int datasync);
20    int (*fasync) (int, struct file *, int);
21    int (*lock) (struct file *, int, struct file_lock *);
22    unsigned long (*get_unmapped_area)(struct file *, unsigned long, unsigned long,
23    unsigned long, unsigned long);
24    int (*check_flags)(int);
25    int (*flock) (struct file *, int, struct file_lock *);
26    ssize_t (*splice_write)(struct pipe_inode_info *, struct file *, loff_t *, size_t,
27    unsigned int);
28    ssize_t (*splice_read)(struct file *, loff_t *, struct pipe_inode_info *, size_t,
29    unsigned int);
30    void (*splice_eof)(struct file *file);
31    int (*setlease)(struct file *, int, struct file_lock **, void **);
32    long (*fallocate)(struct file *file, int mode, loff_t offset,
33    loff_t len);
34    void (*show_fdinfo)(struct seq_file *m, struct file *f);
35    #ifndef CONFIG_MMU
36    unsigned (*mmap_capabilities)(struct file *);
37    #endif
38    ssize_t (*copy_file_range)(struct file *, loff_t, struct file *,
39    loff_t, size_t, unsigned int);
40    loff_t (*remap_file_range)(struct file *file_in, loff_t pos_in,
41    struct file *file_out, loff_t pos_out,
42    loff_t len, unsigned int remap_flags);
43    int (*fadvise)(struct file *, loff_t, loff_t, int);
44    int (*uring_cmd)(struct io_uring_cmd *ioucmd, unsigned int issue_flags);
45    int (*uring_cmd_iopoll)(struct io_uring_cmd *, struct io_comp_batch *,
46    unsigned int poll_flags);
47 } __randomize_layout;

```

```

1 struct seq_file {
2     char *buf;
3     size_t size;
4     size_t from;
5     size_t count;
6     size_t pad_until;
7     loff_t index;
8     loff_t read_pos;
9     struct mutex lock;
10    const struct seq_operations *op;
11    int poll_event;
12    const struct file *file;
13    void *private;
14 };

```

```

1 struct seq_operations {
2     void * (*start) (struct seq_file *m, loff_t *pos);
3     void (*stop) (struct seq_file *m, void *v);
4     void * (*next) (struct seq_file *m, void *v, loff_t *pos);
5     int (*show) (struct seq_file *m, void *v);
6 };

```

```

1 int single_open(struct file *file, int (*show)(struct seq_file *, void *),
2 void *data)
3 {
4     struct seq_operations *op = kmalloc(sizeof(*op), GFP_KERNEL_ACCOUNT);
5     int res = -ENOMEM;
6
7     if (op) {
8         op->start = single_start;
9         op->next = single_next;
10        op->stop = single_stop;
11        op->show = show;
12        res = seq_open(file, op);
13        if (!res)
14            ((struct seq_file *)file->private_data)->private = data;
15        else
16            kfree(op);
17    }
18    return res;
19 }
20 EXPORT_SYMBOL(single_open);

```

```

1 struct kmem_cache {
2     #ifndef CONFIG_SLUB_TINY
3     struct kmem_cache_cpu __percpu *cpu_slab;
4     #endif
5     /* Used for retrieving partial slabs, etc. */
6     slab_flags_t flags;
7     unsigned long min_partial;
8     unsigned int size; /* Object size including metadata */
9     unsigned int object_size; /* Object size without metadata */
10    struct reciprocal_value reciprocal_size;
11    unsigned int offset; /* Free pointer offset */
12    #ifdef CONFIG_SLUB_CPU_PARTIAL
13    /* Number of per cpu partial objects to keep around */
14    unsigned int cpu_partial;
15    /* Number of per cpu partial slabs to keep around */
16    unsigned int cpu_partial_slabs;
17    #endif
18    struct kmem_cache_order_objects oo; /* Allocation and freeing of slabs */
19    struct kmem_cache_order_objects min;
20    gfp_t allocflags; /* gfp flags to use on each alloc */
21    int refcount; /* Refcount for slab cache destroy */
22    void (*ctor)(void *object); /* Object constructor */
23    unsigned int inuse; /* Offset to metadata */
24    unsigned int align; /* Alignment */
25    unsigned int red_left_pad; /* Left redzone padding size */
26    const char *name; /* Name (only for display!) */
27    struct list_head list; /* List of slab caches */
28    #ifdef CONFIG_SYSFS
29    struct kobject kobj; /* For sysfs */
30    #endif
31    #ifdef CONFIG_SLAB_FREELIST_HARDENED
32    unsigned long random;
33    #endif
34    #ifdef CONFIG_NUMA
35    /** Defragmentation by allocating from a remote node. */
36    unsigned int remote_node_defrag_ratio;
37    #endif
38    #ifdef CONFIG_SLAB_FREELIST_RANDOM
39    unsigned int *random_seq;
40    #endif
41    #ifdef CONFIG_KASAN_GENERIC
42    struct kasan_cache kasan_info;
43    #endif
44    #ifdef CONFIG_HARDENED_USERCOPY
45    unsigned int useroffset; /* Usercopy region offset */
46    unsigned int usersize; /* Usercopy region size */
47    #endif
48    struct kmem_cache_node *node[MAX_NUMNODES];
49 };

```

```

1 struct irqaction {
2     irq_handler_t      handler;
3     void              *dev_id;
4     void __percpu      *percpu_dev_id;
5     struct irqaction   *next;
6     irq_handler_t      thread_fn;
7     struct task_struct *thread;
8     struct irqaction   *secondary;
9     unsigned int       irq;
10    unsigned int       flags;
11    unsigned long       thread_flags;
12    unsigned long       thread_mask;
13    const char          *name;
14    struct proc_dir_entry *dir;
15 } ____cacheline_internodealigned_in_smp;

```

```

1 struct softirq_action {
2     void (*action)(struct softirq_action *);
3 };

```

```

1 static struct softirq_action softirq_vec[NR_SOFTIRQS] __cacheline_aligned_in_smp;
2 const char * const softirq_to_name[NR_SOFTIRQS] = {
3     "HI", "TIMER", "NET_TX", "NET_RX", "BLOCK", "IRQ_POLL",
4     "TASKLET", "SCHED", "HRTIMER", "RCU"
5 };

```

```

1 struct tasklet_struct
2 {
3     struct tasklet_struct *next;
4     unsigned long state;
5     atomic_t count;
6     bool use_callback;
7     union {
8         void (*func)(unsigned long data);
9         void (*callback)(struct tasklet_struct *t);
10    };
11    unsigned long data;
12 };

```



```

1 struct workqueue_struct {
2     struct list_head    pwqs;           /* WR: all pwqs of this wq */
3     struct list_head    list;           /* PR: list of all workqueues */
4
5     struct mutex         mutex;          /* protects this wq */
6     int                 work_color;      /* WQ: current work color */
7     int                 flush_color;     /* WQ: current flush color */
8     atomic_t             nr_pwqs_to_flush; /* flush in progress */
9     struct wq_flusher    *first_flusher; /* WQ: first flusher */
10    struct list_head     flusher_queue;  /* WQ: flush waiters */
11    struct list_head     flusher_overflow; /* WQ: flush overflow list */
12
13    struct list_head     maydays;        /* MD: pwqs requesting rescue */
14    struct worker         *rescuer;       /* MD: rescue worker */
15
16    int                 nr_drainers;      /* WQ: drain in progress */
17    int                 saved_max_active; /* WQ: saved pwq max_active */
18
19    struct workqueue_attrs *unbound_attrs; /* PW: only for unbound wqs */
20    struct pool_workqueue *dfl_pwq;      /* PW: only for unbound wqs */
21
22    #ifdef CONFIG_SYSFS
23    struct wq_device      *wq_dev;       /* I: for sysfs interface */
24    #endif
25    #ifdef CONFIG_LOCKDEP
26    char                  *lock_name;
27    struct lock_class_key key;
28    struct lockdep_map    lockdep_map;
29    #endif
30    char                  name[WQ_NAME_LEN]; /* I: workqueue name */
31
32    /*
33     * Destruction of workqueue_struct is RCU protected to allow walking
34     * the workqueues list without grabbing wq_pool_mutex.
35     * This is used to dump all workqueues from sysrq.
36     */
37    struct rcu_head        rcu;
38
39    /* hot fields used during command issue, aligned to cacheline */
40    unsigned int          flags ____cacheline_aligned; /* WQ: WQ_* flags */
41    struct pool_workqueue __percpu __rcu **cpu_pwq; /* I: per-cpu pwqs */
42 };

```

```

1 struct work_struct {
2     atomic_long_t data;
3     struct list_head entry;
4     work_func_t func;
5     #ifdef CONFIG_LOCKDEP
6     struct lockdep_map lockdep_map;
7     #endif
8 };

```

```

1 struct device {
2     struct kobject kobj;
3     struct device      *parent;
4     struct device_private *p;
5     const char          *init_name; /* initial name of the device */
6     const struct device_type *type;
7     const struct bus_type *bus; /* type of bus device is on */
8     struct device_driver *driver; /* which driver has allocated this device */
9     void *platform_data; /* Platform specific data, device
10    core doesn't touch it */
11    void *driver_data; /* Driver data, set and get with
12    dev_set_drvdata/dev_get_drvdata */
13    struct mutex mutex; /* mutex to synchronize calls to * its driver. */
14    struct dev_links_info links;
15    struct dev_pm_info power;
16    struct dev_pm_domain *pm_domain;
17    #ifdef CONFIG_ENERGY_MODEL
18    struct em_perf_domain *em_pd;
19    #endif
20    #ifdef CONFIG_PINCTRL
21    struct dev_pin_info *pins;
22    #endif
23    struct dev_msi_info msi;
24    #ifdef CONFIG_DMA_OPS
25    const struct dma_map_ops *dma_ops;
26    #endif
27    u64 *dma_mask; /* dma mask (if dma'able device) */
28    u64 coherent_dma_mask; /* Like dma_mask, but for
29    alloc_coherent mappings as
30    not all hardware supports
31    64 bit addresses for consistent
32    allocations such descriptors. */
33    u64 bus_dma_limit; /* upstream dma constraint */
34    const struct bus_dma_region *dma_range_map;
35    struct device_dma_parameters *dma_parms;
36    struct list_head dma_pools; /* dma pools (if dma'ble) */
37    #ifdef CONFIG_DMA_DECLARE_COHERENT
38    struct dma_coherent_mem *dma_mem; /* internal for coherent mem override */
39    #endif
40    #ifdef CONFIG_DMA_CMA
41    struct cma *cma_area; /* contiguous memory area for dma allocations */
42    #endif
43    #ifdef CONFIG_SWIOTLB
44    struct io_tlb_mem *dma_io_tlb_mem;
45    #endif
46    #ifdef CONFIG_SWIOTLB_DYNAMIC
47    struct list_head dma_io_tlb_pools;
48    spinlock_t dma_io_tlb_lock;
49    bool dma_uses_io_tlb;
50    #endif
51    struct dev_archdata archdata; /* arch specific additions */
52    struct device_node *of_node; /* associated device tree node */
53    struct fwnode_handle *fwnode; /* firmware device node */
54    #ifdef CONFIG_NUMA
55    int numa_node; /* NUMA node this device is close to */
56    #endif
57    dev_t devt; /* dev_t, creates the sysfs "dev" */
58    u32 id; /* device instance */
59    spinlock_t devres_lock;
60    struct list_head devres_head;
61    const struct class *class;
62    const struct attribute_group **groups; /* optional groups */
63    void (*release)(struct device *dev);
64    struct iommu_group *iommu_group;
65    struct dev_iommu *iommu;
66    struct device_physical_location *physical_location;
67    enum device_removable removable;
68    bool offline_disabled:1;
69    bool offline:1;
70    bool of_node_reused:1;

```

```

71     bool                state_synced:1;
72     bool                can_match:1;
73     #if defined(CONFIG_ARCH_HAS_SYNC_DMA_FOR_DEVICE) || \
74     defined(CONFIG_ARCH_HAS_SYNC_DMA_FOR_CPU) || \
75     defined(CONFIG_ARCH_HAS_SYNC_DMA_FOR_CPU_ALL)
76     bool                dma_coherent:1;
77     #endif
78     #ifdef CONFIG_DMA_OPS_BYPASS
79     bool                dma_ops_bypass : 1;
80     #endif
81 };

```

```

1  struct device_driver {
2      const char          *name;
3      const struct bus_type *bus;
4      struct module       *owner;
5      const char          *mod_name;          /* used for built-in modules */
6      bool suppress_bind_attrs;              /* disables bind/unbind via sysfs */
7      enum probe_type probe_type;
8      const struct of_device_id *of_match_table;
9      const struct acpi_device_id *acpi_match_table;
10     int (*probe) (struct device *dev);
11     void (*sync_state) (struct device *dev);
12     int (*remove) (struct device *dev);
13     void (*shutdown) (struct device *dev);
14     int (*suspend) (struct device *dev, pm_message_t state);
15     int (*resume) (struct device *dev);
16     const struct attribute_group **groups;
17     const struct attribute_group **dev_groups;
18     const struct dev_pm_ops *pm;
19     void (*coredump) (struct device *dev);
20     struct driver_private *p;
21 };

```

```

1 struct usb_interface {
2     /* array of alternate settings for this interface,
3      * stored in no particular order */
4     struct usb_host_interface *altsetting;
5
6     struct usb_host_interface *cur_altsetting;    /* the currently
7      * active alternate setting */
8     unsigned num_altsetting;    /* number of alternate settings */
9
10    /* If there is an interface association descriptor then it will list
11     * the associated interfaces */
12    struct usb_interface_assoc_descriptor *intf_assoc;
13
14    int minor;    /* minor number this interface is
15     * bound to */
16    enum usb_interface_condition condition;    /* state of binding */
17    unsigned sysfs_files_created:1; /* the sysfs attributes exist */
18    unsigned ep_devs_created:1;    /* endpoint "devices" exist */
19    unsigned unregistering:1;    /* unregistration is in progress */
20    unsigned needs_remote_wakeup:1; /* driver requires remote wakeup */
21    unsigned needs_altsetting0:1; /* switch to altsetting 0 is pending */
22    unsigned needs_binding:1;    /* needs delayed unbind/rebind */
23    unsigned resetting_device:1; /* true: bandwidth alloc after reset */
24    unsigned authorized:1;    /* used for interface authorization */
25    enum usb_wireless_status wireless_status;
26    struct work_struct wireless_status_work;
27
28    struct device dev;    /* interface specific device info */
29    struct device *usb_dev;
30    struct work_struct reset_ws;    /* for resets in atomic context */
31 };

```

```

1 struct urb {
2     /* private: usb core and host controller only fields in the urb */
3     struct kref kref;    /* reference count of the URB */
4     int unlinked;    /* unlink error code */
5     void *hcpriv;    /* private data for host controller */
6     atomic_t use_count;    /* concurrent submissions counter */
7     atomic_t reject;    /* submissions will fail */
8     /* public: documented fields in the urb that can be used by drivers */
9     struct list_head urb_list;    /* list head for use by the urb's * current owner */
10    struct list_head anchor_list;    /* the URB may be anchored */
11    struct usb_anchor *anchor;
12    struct usb_device *dev;    /* (in) pointer to associated device */
13    struct usb_host_endpoint *ep;    /* (internal) pointer to endpoint */
14    unsigned int pipe;    /* (in) pipe information */
15    unsigned int stream_id;    /* (in) stream ID */
16    int status;    /* (return) non-ISO status */
17    unsigned int transfer_flags;    /* (in) URB_SHORT_NOT_OK | ... */
18    void *transfer_buffer;    /* (in) associated data buffer */
19    dma_addr_t transfer_dma;    /* (in) dma addr for transfer_buffer */
20    struct scatterlist *sg;    /* (in) scatter gather buffer list */
21    int num_mapped_sgs;    /* (internal) mapped sg entries */
22    int num_sgs;    /* (in) number of entries in the sg list */
23    u32 transfer_buffer_length;    /* (in) data buffer length */
24    u32 actual_length;    /* (return) actual transfer length */
25    unsigned char *setup_packet;    /* (in) setup packet (control only) */
26    dma_addr_t setup_dma;    /* (in) dma addr for setup_packet */
27    int start_frame;    /* (modify) start frame (ISO) */
28    int number_of_packets;    /* (in) number of ISO packets */
29    int interval;    /* (modify) transfer interval * (INT/ISO) */
30    int error_count;    /* (return) number of ISO errors */
31    void *context;    /* (in) context for completion */
32    usb_complete_t complete;    /* (in) completion routine */
33    struct usb_iso_packet_descriptor iso_frame_desc[]; /* (in) ISO ONLY */
34 };

```

```

1 struct usb_driver {
2     const char *name;
3
4     int (*probe) (struct usb_interface *intf,
5     const struct usb_device_id *id);
6
7     void (*disconnect) (struct usb_interface *intf);
8
9     int (*unlocked_ioctl) (struct usb_interface *intf, unsigned int code,
10    void *buf);
11
12    int (*suspend) (struct usb_interface *intf, pm_message_t message);
13    int (*resume) (struct usb_interface *intf);
14    int (*reset_resume)(struct usb_interface *intf);
15
16    int (*pre_reset)(struct usb_interface *intf);
17    int (*post_reset)(struct usb_interface *intf);
18
19    const struct usb_device_id *id_table;
20    const struct attribute_group **dev_groups;
21
22    struct usb_dynids dynids;
23    struct device_driver driver;
24    unsigned int no_dynamic_id:1;
25    unsigned int supports_autosuspend:1;
26    unsigned int disable_hub_initiated_lpm:1;
27    unsigned int soft_unbind:1;
28 };

```

```

1 struct usb_device_driver {
2     const char *name;
3
4     bool (*match) (struct usb_device *udev);
5     int (*probe) (struct usb_device *udev);
6     void (*disconnect) (struct usb_device *udev);
7
8     int (*suspend) (struct usb_device *udev, pm_message_t message);
9     int (*resume) (struct usb_device *udev, pm_message_t message);
10
11    int (*choose_configuration) (struct usb_device *udev);
12
13    const struct attribute_group **dev_groups;
14    struct device_driver driver;
15    const struct usb_device_id *id_table;
16    unsigned int supports_autosuspend:1;
17    unsigned int generic_subclass:1;
18 };

```

```

1 struct task_struct {
2     #ifdef CONFIG_THREAD_INFO_IN_TASK
3     struct thread_info thread_info;
4     #endif
5     unsigned int    __state;
6     /* saved state for "spinlock sleepers" */
7     unsigned int    saved_state;
8     /*
9     * This begins the randomizable portion of task_struct. Only
10    * scheduling-critical items should be added above here.
11    */
12    randomized_struct_fields_start
13    void*stack;
14    refcount_t       usage;
15    /* Per task flags (PF_*), defined further below: */
16    unsigned int     flags;
17    unsigned int     ptrace;
18    #ifdef CONFIG_MEM_ALLOC_PROFILING
19    struct alloc_tag *alloc_tag;
20    #endif
21    #ifdef CONFIG_SMP
22    int on_cpu;
23    struct __call_single_node    wake_entry;
24    unsigned int     wakee_flips;
25    unsigned long    wakee_flip_decay_ts;
26    struct task_struct *last_wakee;
27    /*
28    * recent_used_cpu is initially set as the last CPU used by a task
29    * that wakes affine another task. Waker/wakee relationships can
30    * push tasks around a CPU where each wakeup moves to the next one.
31    * Tracking a recently used CPU allows a quick search for a recently
32    * used CPU that may be idle.
33    */
34    int recent_used_cpu;
35    int wake_cpu;
36    #endif
37    int on_rq;
38    int prio;
39    int static_prio;
40    int normal_prio;
41    unsigned int     rt_priority;
42    struct sched_entity se;
43    struct sched_rt_entity rt;
44    struct sched_dl_entity dl;
45    struct sched_dl_entity *dl_server;
46    #ifdef CONFIG_SCHED_CLASS_EXT
47    struct sched_ext_entity scx;
48    #endif
49    const struct sched_class    *sched_class;
50    #ifdef CONFIG_SCHED_CORE
51    struct rb_node    core_node;
52    unsigned long     core_cookie;
53    unsigned int      core_occupation;
54    #endif
55    #ifdef CONFIG_CGROUP_SCHED
56    struct task_group *sched_task_group;
57    #endif
58    #ifdef CONFIG_UCLAMP_TASK
59    /*
60    * Clamp values requested for a scheduling entity.
61    * Must be updated with task_rq_lock() held.
62    */
63    struct uclamp_se uclamp_req[UCLAMP_CNT];
64    /*
65    * Effective clamp values used for a scheduling entity.
66    * Must be updated with task_rq_lock() held.
67    */
68    struct uclamp_se uclamp[UCLAMP_CNT];
69    #endif
70    struct sched_statistics    stats;

```

```

71 #ifdef CONFIG_PREEMPT_NOTIFIERS
72 /* List of struct preempt_notifier: */
73 struct hlist_head preempt_notifiers;
74 #endif
75 #ifdef CONFIG_BLK_DEV_IO_TRACE
76 unsigned int    btrace_seq;
77 #endif
78 unsigned int    policy;
79 unsigned long   max_allowed_capacity;
80 int nr_cpus_allowed;
81 const cpumask_t *cpus_ptr;
82 cpumask_t       *user_cpus_ptr;
83 cpumask_t       cpus_mask;
84 void *migration_pending;
85 #ifdef CONFIG_SMP
86 unsigned short  migration_disabled;
87 #endif
88 unsigned short  migration_flags;
89 #ifdef CONFIG_PREEMPT_RCU
90 int rcu_read_lock_nesting;
91 union rcu_special rcu_read_unlock_special;
92 struct list_head rcu_node_entry;
93 struct rcu_node *rcu_blocked_node;
94 #endif /* #ifdef CONFIG_PREEMPT_RCU */
95 #ifdef CONFIG_TASKS_RCU
96 unsigned long   rcu_tasks_nvcsw;
97 u8rcu_tasks_holdout;
98 u8rcu_tasks_idx;
99 int rcu_tasks_idle_cpu;
100 struct list_headrcu_tasks_holdout_list;
101 int rcu_tasks_exit_cpu;
102 struct list_head rcu_tasks_exit_list;
103 #endif /* #ifdef CONFIG_TASKS_RCU */
104 #ifdef CONFIG_TASKS_TRACE_RCU
105 int trc_reader_nesting;
106 int trc_ipi_to_cpu;
107 union rcu_special trc_reader_special;
108 struct list_head trc_holdout_list;
109 struct list_head trc_blkdn_node;
110 int trc_blkdn_cpu;
111 #endif /* #ifdef CONFIG_TASKS_TRACE_RCU */
112 struct sched_info sched_info;
113 struct list_head tasks;
114 #ifdef CONFIG_SMP
115 struct plist_node pushable_tasks;
116 struct rb_node   pushable_dl_tasks;
117 #endif
118 struct mm_struct *mm;
119 struct mm_struct *active_mm;
120 struct address_space *faults_disabled_mapping;
121 int exit_state;
122 int exit_code;
123 int exit_signal;
124 /* The signal sent when the parent dies: */
125 int pdeath_signal;
126 /* JOBCTL_*, siglock protected: */
127 unsigned long   jobctl;
128 /* Used for emulating ABI behavior of previous Linux versions: */
129 unsigned int    personality;
130 /* Scheduler bits, serialized by scheduler locks: */
131 unsigned        sched_reset_on_fork:1;
132 unsigned        sched_contributes_to_load:1;
133 unsigned        sched_migrated:1;
134 unsigned        sched_task_hot:1;
135 /* Force alignment to the next boundary: */
136 unsigned        :0;
137 /* Unserialized, strictly 'current' */
138 /*
139 * This field must not be in the scheduler word above due to wakelist
140 * queueing no longer being serialized by p->on_cpu. However:
141 * p->XXX = X;    ttwu()

```

```

142 * schedule()          if (p->on_rq && ...) // false
143 *   smp_mb__after_spinlock();          if (smp_load_acquire(&p->on_cpu) && //true
144 *   deactivate_task()          ttwu_queue_wakelist())
145 *   p->on_rq = 0;          p->sched_remote_wakeup = Y;
146 * guarantees all stores of 'current' are visible before
147 * ->sched_remote_wakeup gets used, so it can be in this word.
148 */
149 unsigned          sched_remote_wakeup:1;
150 #ifdef CONFIG_RT_MUTEXES
151 unsigned          sched_rt_mutex:1;
152 #endif
153 /* Bit to tell TOMOYO we're in execve(): */
154 unsigned          in_execve:1;
155 unsigned          in_iowait:1;
156 #ifndef TIF_RESTORE_SIGMASK
157 unsigned          restore_sigmask:1;
158 #endif
159 #ifdef CONFIG_MEMCG_V1
160 unsigned          in_user_fault:1;
161 #endif
162 #ifdef CONFIG_LRU_GEN
163 /* whether the LRU algorithm may apply to this access */
164 unsigned          in_lru_fault:1;
165 #endif
166 #ifdef CONFIG_COMPAT_BRK
167 unsigned          brk_randomized:1;
168 #endif
169 #ifdef CONFIG_CGROUPS
170 /* disallow userland-initiated cgroup migration */
171 unsigned          no_cgroup_migration:1;
172 /* task is frozen/stopped (used by the cgroup freezer) */
173 unsigned          frozen:1;
174 #endif
175 #ifdef CONFIG_BLK_CGROUP
176 unsigned          use_memdelay:1;
177 #endif
178 #ifdef CONFIG_PSI
179 /* Stalled due to lack of memory */
180 unsigned          in_memstall:1;
181 #endif
182 #ifdef CONFIG_PAGE_OWNER
183 /* Used by page_owner=on to detect recursion in page tracking. */
184 unsigned          in_page_owner:1;
185 #endif
186 #ifdef CONFIG_EVENTFD
187 /* Recursion prevention for eventfd_signal() */
188 unsigned          in_eventfd:1;
189 #endif
190 #ifdef CONFIG_ARCH_HAS_CPU_PASID
191 unsigned          pasid_activated:1;
192 #endif
193 #ifdef CONFIG_X86_BUS_LOCK_DETECT
194 unsigned          reported_split_lock:1;
195 #endif
196 #ifdef CONFIG_TASK_DELAY_ACCT
197 /* delay due to memory thrashing */
198 unsigned          in_thrashing:1;
199 #endif
200 #ifdef CONFIG_PREEMPT_RT
201 struct netdev_xmitnet_xmit;
202 #endif
203 unsigned long      atomic_flags; /* Flags requiring atomic access. */
204 struct restart_blockrestart_block;
205 pid_t pid;
206 pid_t tgid;
207 #ifdef CONFIG_STACKPROTECTOR
208 /* Canary value for the -fstack-protector GCC feature: */
209 unsigned long      stack_canary;
210 #endif
211 /*
212 * Point ers to the (original) parent process, youngest child, younger sibling,

```



```

213 * older sibling, respectively. (p->father can be replaced with
214 * p->real_parent->pid)
215 */
216 /* Real parent process: */
217 struct task_struct __rcu    *real_parent;
218 /* Recipient of SIGCHLD, wait4() reports: */
219 struct task_struct __rcu    *parent;
220 /*
221 * Children/sibling form the list of natural children:
222 */
223 struct list_head children;
224 struct list_head sibling;
225 struct task_struct *group_leader;
226 /*
227 * 'ptraced' is the list of tasks this task is using ptrace() on.
228 * This includes both natural children and PTRACE_ATTACH targets.
229 * 'ptrace_entry' is this task's link on the p->parent->ptraced list.
230 */
231 struct list_head ptraced;
232 struct list_head ptrace_entry;
233 /* PID/PID hash table linkage. */
234 struct pid        *thread_pid;
235 struct hlist_node pid_links[PIDTYPE_MAX];
236 struct list_head thread_node;
237 struct completion *vfork_done;
238 /* CLONE_CHILD_SETTID: */
239 int __user        *set_child_tid;
240 /* CLONE_CHILD_CLEARTID: */
241 int __user        *clear_child_tid;
242 /* PF_KTHREAD | PF_IO_WORKER */
243 void *worker_private;
244 u64 utime;
245 u64 stime;
246 #ifdef CONFIG_ARCH_HAS_SCALED_CPUTIME
247 u64 utimescaled;
248 u64 stimescaled;
249 #endif
250 u64 gtime;
251 struct prev_cpu timeprev_cputime;
252 #ifdef CONFIG_VIRT_CPU_ACCOUNTING_GEN
253 struct vtime      vtime;
254 #endif
255 #ifdef CONFIG_NO_HZ_FULL
256 atomic_t          tick_dep_mask;
257 #endif
258 /* Context switch counts: */
259 unsigned long     nvcs;
260 unsigned long     nivcs;
261 /* Monotonic time in nsecs: */
262 u64 start_time;
263 /* Boot based time in nsecs: */
264 u64 start_boottime;
265 /* MM fault and swap info: this can arguably be seen as either mm-specific or
thread-specific: */
266 unsigned long     min_flt;
267 unsigned long     maj_flt;
268 /* Empty if CONFIG_POSIX_CPUTIMERS=n */
269 struct posix_cputimers posix_cputimers;
270 #ifdef CONFIG_POSIX_CPU_TIMERS_TASK_WORK
271 struct posix_cputimers_work posix_cputimers_work;
272 #endif
273 /* Process credentials: */
274 /* Tracer's credentials at attach: */
275 const struct cred __rcu *ptracer_cred;
276 /* Objective and real subjective task credentials (COW): */
277 const struct cred __rcu *real_cred;
278 /* Effective (overridable) subjective task credentials (COW): */
279 const struct cred __rcu *cred;
280 #ifdef CONFIG_KEYS
281 /* Cached requested key. */
282 struct key        *cached_requested_key;

```

```

283 #endif
284 /*
285  * executable name, excluding path.
286  * - normally initialized begin_new_exec()
287  * - set it with set_task_comm()
288  * - strncpy_pad() to ensure it is always NUL-terminated and zero-padded
289  * - task_lock() to ensure the operation is atomic and the name is fully updated.
290  */
291 char comm[TASK_COMM_LEN];
292 struct nameidata*nameidata;
293 #ifdef CONFIG_SYSVIPC
294 struct sysv_sem sysvsem;
295 struct sysv_shm sysvshm;
296 #endif
297 #ifdef CONFIG_DETECT_HUNG_TASK
298 unsigned long last_switch_count;
299 unsigned long last_switch_time;
300 #endif
301 /* Filesystem information: */
302 struct fs_struct *fs;
303 /* Open file information: */
304 struct files_struct *files;
305 #ifdef CONFIG_IO_URING
306 struct io_uring_task*io_uring;
307 #endif
308 /* Namespaces: */
309 struct nsproxy *nsproxy;
310 /* Signal handlers: */
311 struct signal_struct *signal;
312 struct sighand_struct __rcu*sighand;
313 sigset_t blocked;
314 sigset_t real_blocked;
315 /* Restored if set_restore_sigmask() was used: */
316 sigset_t saved_sigmask;
317 struct sigpendingpending;
318 unsigned long sas_ss_sp;
319 size_t sas_ss_size;
320 unsigned int sas_ss_flags;
321 struct callback_head *task_works;
322 #ifdef CONFIG_AUDIT
323 #ifdef CONFIG_AUDITSYSCALL
324 struct audit_context *audit_context;
325 #endif
326 kuid_t loginuid;
327 unsigned int sessionid;
328 #endif
329 struct seccomp seccomp;
330 struct syscall_user_dispatch syscall_dispatch;
331 /* Thread group tracking: */
332 u64 parent_exec_id;
333 u64 self_exec_id;
334 /* Protection against (de-)allocation: mm, files, fs, tty, keyrings, mems_allowed,
mempolicy: */
335 spinlock_t alloc_lock;
336 /* Protection of the PI data structures: */
337 raw_spinlock_t pi_lock;
338 struct wake_q_nodewake_q;
339 #ifdef CONFIG_RT_MUTEXES
340 /* PI waiters blocked on a rt_mutex held by this task: */
341 struct rb_root_cached pi_waiters;
342 /* Updated under owner's pi_lock and rq lock */
343 struct task_struct *pi_top_task;
344 /* Deadlock detection and priority inheritance handling: */
345 struct rt_mutex_waiter *pi_blocked_on;
346 #endif
347 #ifdef CONFIG_DEBUG_MUTEXES
348 /* Mutex deadlock detection: */
349 struct mutex_waiter *blocked_on;
350 #endif
351 #ifdef CONFIG_DETECT_HUNG_TASK_BLOCKER
352 struct mutex *blocker_mutex;

```

```

353 #endif
354 #ifdef CONFIG_DEBUG_ATOMIC_SLEEP
355 int non_block_count;
356 #endif
357 #ifdef CONFIG_TRACE_IRQFLAGS
358 struct irqtrace_eventsirqtrace;
359 unsigned int    hardirq_threaded;
360 u64 hardirq_chain_key;
361 int softirqs_enabled;
362 int softirq_context;
363 int irq_config;
364 #endif
365 #ifdef CONFIG_PREEMPT_RT
366 int softirq_disable_cnt;
367 #endif
368 #ifdef CONFIG_LOCKDEP
369 # define MAX_LOCK_DEPTH 48UL
370 u64 curr_chain_key;
371 int lockdep_depth;
372 unsigned int    lockdep_recursion;
373 struct held_lockheld_locks[MAX_LOCK_DEPTH];
374 #endif
375 #if defined(CONFIG_UBSAN) && !defined(CONFIG_UBSAN_TRAP)
376 unsigned int    in_ubsan;
377 #endif
378 /* Journalling filesystem info: */
379 void * journal_info;
380 /* Stacked block device info: */
381 struct bio_list *bio_list;
382 /* Stack plugging: */
383 struct blk_plug *plug;
384 /* VM state: */
385 struct reclaim_state*reclaim_state;
386 struct io_context*io_context;
387 #ifdef CONFIG_COMPACTION
388 struct capture_control*capture_control;
389 #endif
390 /* Ptrace state: */
391 unsigned long    ptrace_message;
392 kernel_siginfo_t*last_siginfo;
393 struct task_io_accounting    ioac;
394 #ifdef CONFIG_PSI
395 /* Pressure stall state */
396 unsigned int    psi_flags;
397 #endif
398 #ifdef CONFIG_TASK_XACCT
399 /* Accumulated RSS usage: */
400 u64 acct_rss_mem1;
401 /* Accumulated virtual memory usage: */
402 u64 acct_vm_mem1;
403 /* stime + utime since last update: */
404 u64 acct_timexpd;
405 #endif
406 #ifdef CONFIG_CPUSETS
407 /* Protected by ->alloc_lock: */
408 nodemask_t    mems_allowed;
409 /* Sequence number to catch updates: */
410 seqcount_spinlock_tmems_allowed_seq;
411 int cpuset_mem_spread_rotor;
412 #endif
413 #ifdef CONFIG_CGROUPS
414 /* Control Group info protected by css_set_lock: */
415 struct css_set __rcu *cgroups;
416 /* cg_list protected by css_set_lock and tsk->alloc_lock: */
417 struct list_head cg_list;
418 #endif
419 #ifdef CONFIG_X86_CPU_RESCTRL
420 u32 closid;
421 u32 rmid;
422 #endif
423 #ifdef CONFIG_FUTEX

```

```

424 struct robust_list_head __user *robust_list;
425 #ifdef CONFIG_COMPAT
426 struct compat_robust_list_head __user *compat_robust_list;
427 #endif
428 struct list_head pi_state_list;
429 struct futex_pi_state *pi_state_cache;
430 struct mutex futex_exit_mutex;
431 unsigned int futex_state;
432 #endif
433 #ifdef CONFIG_PERF_EVENTS
434 u8perf_recursion[PERF_NR_CONTEXTS];
435 struct perf_event_context *perf_event_ctxp;
436 struct mutex perf_event_mutex;
437 struct list_headperf_event_list;
438 struct perf_ctx_data __rcu *perf_ctx_data;
439 #endif
440 #ifdef CONFIG_DEBUG_PREEMPT
441 unsigned long preempt_disable_ip;
442 #endif
443 #ifdef CONFIG_NUMA
444 /* Protected by alloc_lock: */
445 struct mempolicy*mempolicy;
446 short il_prev;
447 u8 il_weight;
448 short pref_node_fork;
449 #endif
450 #ifdef CONFIG_NUMA_BALANCING
451 int numa_scan_seq;
452 unsigned int numa_scan_period;
453 unsigned int numa_scan_period_max;
454 int numa_preferred_nid;
455 unsigned long numa_migrate_retry;
456 /* Migration stamp: */
457 u64 node_stamp;
458 u64 last_task_numa_placement;
459 u64 last_sum_exec_runtime;
460 struct callback_headnuma_work;
461 /*
462  * This point er is only modified for current in syscall and
463  * pagefault context (and for tasks being destroyed), so it can be read
464  * from any of the following contexts:
465  * - RCU read-side critical section
466  * - current->numa_group from everywhere
467  * - task's runqueue locked, task not running
468  */
469 struct numa_group __rcu* numa_group;
470 /*
471  * numa_faults is an array split into four regions:
472  * faults_memory, faults_cpu, faults_memory_buffer, faults_cpu_buffer
473  * in this precise order.
474  * faults_memory: Exponential decaying average of faults on a per-node
475  * basis. Scheduling placement decisions are made based on these
476  * counts. The values remain static for the duration of a PTE scan.
477  * faults_cpu: Track the nodes the process was running on when a NUMA
478  * hinting fault was incurred.
479  * faults_memory_buffer and faults_cpu_buffer: Record faults per node
480  * during the current scan window. When the scan completes, the counts
481  * in faults_memory and faults_cpu decay and these values are copied.
482  */
483 unsigned long *numa_faults;
484 unsigned long total_numa_faults;
485 /*
486  * numa_faults_locality tracks if faults recorded during the last
487  * scan window were remote/local or failed to migrate. The task scan
488  * period is adapted based on the locality of the faults with different
489  * weights depending on whether they were shared or private faults
490  */
491 unsigned long numa_faults_locality[3];
492 unsigned long numa_pages_migrated;
493 #endif /* CONFIG_NUMA_BALANCING */
494 #ifdef CONFIG_RSEQ

```

```

495 struct rseq __user *rseq;
496 u32 rseq_len;
497 u32 rseq_sig;
498 /*
499  * RmW on rseq_event_mask must be performed atomically
500  * with respect to preemption.
501  */
502 unsigned long rseq_event_mask;
503 #ifdef CONFIG_DEBUG_RSEQ
504 /*
505  * This is a place holder to save a copy of the rseq fields for
506  * validation of read-only fields. The struct rseq has a
507  * variable-length array at the end, so it cannot be used
508  * directly. Reserve a size large enough for the known fields.
509  */
510 char rseq_fields[sizeof(struct rseq)];
511 #endif
512 #endif
513 #ifdef CONFIG_SCHED_MM_CID
514 int mm_cid; /* Current cid in mm */
515 int last_mm_cid; /* Most recent cid in mm */
516 int migrate_from_cpu;
517 int mm_cid_active; /* Whether cid bitmap is active */
518 struct callback_headcid_work;
519 #endif
520 struct tlbflush_unmap_batch tlb_ubb;
521 /* Cache last used pipe for splice(): */
522 struct pipe_inode_info*splice_pipe;
523 struct page_fragtask_frag;
524 #ifdef CONFIG_TASK_DELAY_ACCT
525 struct task_delay_info*delays;
526 #endif
527 #ifdef CONFIG_FAULT_INJECTION
528 int make_it_fail;
529 unsigned int fail_nth;
530 #endif
531 /*
532  * When (nr_dirtied >= nr_dirtied_pause), it's time to call
533  * balance_dirty_pages() for a dirty throttling pause:
534  */
535 int nr_dirtied;
536 int nr_dirtied_pause;
537 /* Start of a write-and-pause period: */
538 unsigned long dirty_paused_when;
539 #ifdef CONFIG_LATENCYTOP
540 int latency_record_count;
541 struct latency_recordlatency_record[LT_SAVECOUNT];
542 #endif
543 /*
544  * Time slack values; these are used to round up poll() and
545  * select() etc timeout values. These are in nanoseconds.
546  */
547 u64 timer_slack_ns;
548 u64 default_timer_slack_ns;
549 #if defined(CONFIG_KASAN_GENERIC) || defined(CONFIG_KASAN_SW_TAGS)
550 unsigned int kasan_depth;
551 #endif
552 #ifdef CONFIG_KCSAN
553 struct kcsan_ctxkcsan_ctx;
554 #ifdef CONFIG_TRACE_IRQFLAGS
555 struct irqtrace_eventskcsan_save_irqtrace;
556 #endif
557 #ifdef CONFIG_KCSAN_WEAK_MEMORY
558 int kcsan_stack_depth;
559 #endif
560 #endif
561 #ifdef CONFIG_KMSAN
562 struct kmsan_ctxkmsan_ctx;
563 #endif
564 #if IS_ENABLED(CONFIG_KUNIT)
565 struct kunit *kunit_test;

```

```

566 #endif
567 #ifndef CONFIG_FUNCTION_GRAPH_TRACER
568 /* Index of current stored address in ret_stack: */
569 int curr_ret_stack;
570 int curr_ret_depth;
571 /* Stack of return addresses for return function tracing: */
572 unsigned long *ret_stack;
573 /* Timestamp for last schedule: */
574 unsigned long longftrace_timestamp;
575 unsigned long longftrace_sleeptime;
576 /*
577 * Number of functions that haven't been traced
578 * because of depth overrun:
579 */
580 atomic_t trace_overrun;
581 /* Pause tracing: */
582 atomic_t tracing_graph_pause;
583 #endif
584 #ifdef CONFIG_TRACING
585 /* Bitmask and counter of trace recursion: */
586 unsigned long trace_recursion;
587 #endif /* CONFIG_TRACING */
588 #ifdef CONFIG_KCOV
589 /* See kernel/kcov.c for more details. */
590 /* Coverage collection mode enabled for this task (0 if disabled): */
591 unsigned int kcov_mode;
592 /* Size of the kcov_area: */
593 unsigned int kcov_size;
594 /* Buffer for coverage collection: */
595 void *kcov_area;
596 /* KCOV descriptor wired with this task or NULL: */
597 struct kcov *kcov;
598 /* KCOV common handle for remote coverage collection: */
599 u64 kcov_handle;
600 /* KCOV sequence number: */
601 int kcov_sequence;
602 /* Collect coverage from softirq context: */
603 unsigned int kcov_softirq;
604 #endif
605 #ifdef CONFIG_MEMCG_V1
606 struct mem_cgroup *memcg_in_oom;
607 #endif
608 #ifdef CONFIG_MEMCG
609 /* Number of pages to reclaim on returning to userland: */
610 unsigned int memcg_nr_pages_over_high;
611 /* Used by memcontrol for targeted memcg charge: */
612 struct mem_cgroup*active_memcg;
613 /* Cache for current->cgroups->memcg->objcg lookups: */
614 struct obj_cgroup *objcg;
615 #endif
616 #ifdef CONFIG_BLK_CGROUP
617 struct gendisk *throttle_disk;
618 #endif
619 #ifdef CONFIG_UPROBES
620 struct uprobe_task *utask;
621 #endif
622 #if defined(CONFIG_BCACHE) || defined(CONFIG_BCACHE_MODULE)
623 unsigned int sequential_io;
624 unsigned int sequential_io_avg;
625 #endif
626 struct kmap_ctrlkmap_ctrl;
627 #ifdef CONFIG_DEBUG_ATOMIC_SLEEP
628 unsigned long task_state_change;
629 # ifdef CONFIG_PREEMPT_RT
630 unsigned long saved_state_change;
631 # endif
632 #endif
633 struct rcu_head rcu;
634 refcount_t rcu_users;
635 int pagefault_disabled;
636 #ifdef CONFIG_MMU

```

```

637 struct task_struct *oom_reaper_list;
638 struct timer_list oom_reaper_timer;
639 #endif
640 #ifdef CONFIG_VMAP_STACK
641 struct vm_struct *stack_vm_area;
642 #endif
643 #ifdef CONFIG_THREAD_INFO_IN_TASK
644 /* A live task holds one reference: */
645 refcount_t      stack_refcount;
646 #endif
647 #ifdef CONFIG_LIVEPATCH
648 int patch_state;
649 #endif
650 #ifdef CONFIG_SECURITY
651 /* Used by LSM modules for access restriction: */
652 void * security;
653 #endif
654 #ifdef CONFIG_BPF_SYSCALL
655 /* Used by BPF task local storage */
656 struct bpf_local_storage __rcu *bpf_storage;
657 /* Used for BPF run context */
658 struct bpf_run_ctx*bpf_ctx;
659 #endif
660 /* Used by BPF for per-TASK xdp storage */
661 struct bpf_net_context*bpf_net_context;
662 #ifdef CONFIG_GCC_PLUGIN_STACKLEAK
663 unsigned long lowest_stack;
664 unsigned long prev_lowest_stack;
665 #endif
666 #ifdef CONFIG_X86_MCE
667 void __user *mce_vaddr;
668 __u64 mce_kflags;
669 u64 mce_addr;
670 __u64 mce_ripv : 1,
671 mce_whole_page : 1,
672 __mce_reserved : 62;
673 struct callback_headmce_kill_me;
674 int mce_count;
675 #endif
676 #ifdef CONFIG_KRETPROBES
677 struct llist_head kretprobe_instances;
678 #endif
679 #ifdef CONFIG_RETHOOK
680 struct llist_head rethooks;
681 #endif
682 #ifdef CONFIG_ARCH_HAS_PARANOID_L1D_FLUSH
683 /*
684  * If L1D flush is supported on mm context switch
685  * then we use this callback head to queue kill work
686  * to kill tasks that are not running on SMT disabled
687  * cores
688  */
689 struct callback_headl1d_flush_kill;
690 #endif
691 #ifdef CONFIG_RV
692 /*
693  * Per-task RV monitor. Nowadays fixed in RV_PER_TASK_MONITORS.
694  * If we find justification for more monitors, we can think
695  * about adding more or developing a dynamic method. So far,
696  * none of these are justified.
697  */
698 union rv_task_monitorr[RV_PER_TASK_MONITORS];
699 #endif
700 #ifdef CONFIG_USER_EVENTS
701 struct user_event_mm *user_event_mm;
702 #endif
703 randomized_struct_fields_end
704 /* CPU-specific state of this task: */
705 struct thread_struct thread;
706 };

```